

---

# GR- Router : A Distributed GNU Radio Framework

Tommy Tracy II

Computer Engineering Graduate Student

---

Professor Mircea Stan

HPLP Advisor

Alliant Techsystems Inc. (ATK)

Funding Organization

# Outline

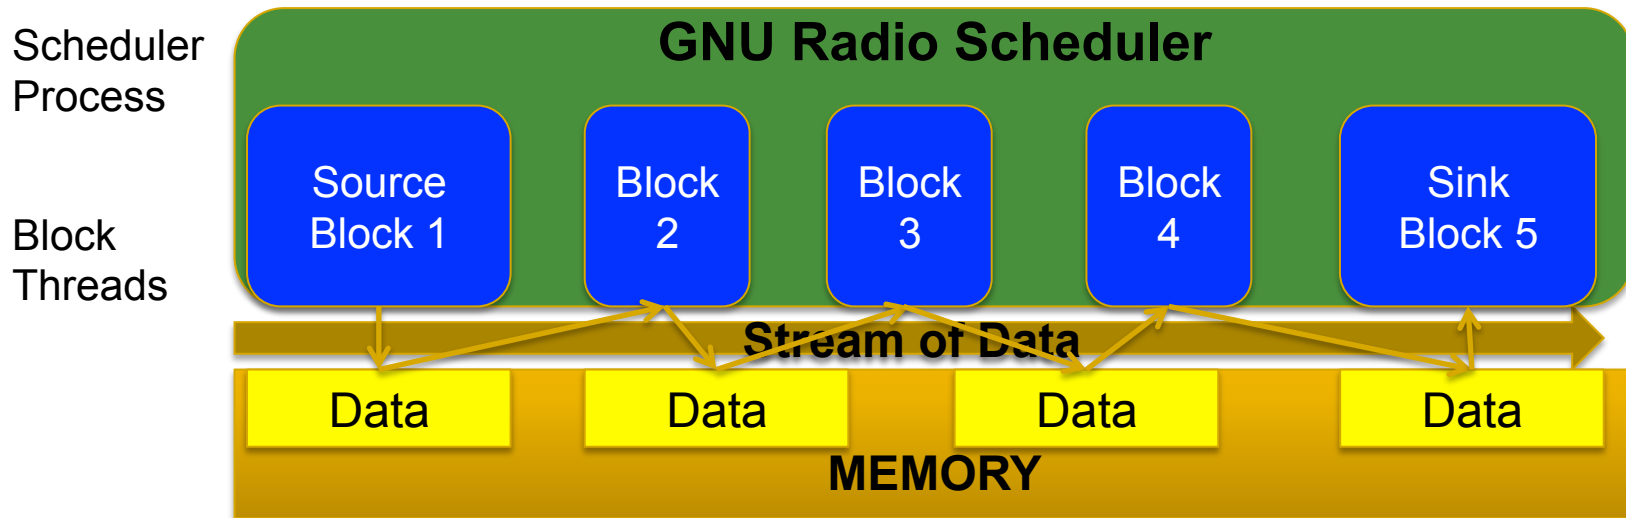
- **Background**
- **Problem Statement** : Limited scalability in high-throughput shared-memory SDR system.
- **Solution** : Distribute the SDR system with GR-Router.
- **Example Implementation** : Test with GR-LDPC.
- **Future Work**



# Background

## ■ GNU Radio

- ❑ Is an open-source software development toolkit.
- ❑ It provides DSP blocks for SDR applications.
- ❑ Each GNU Radio block is spawned as an independent processing thread which uses inter-thread communication mechanisms (via the scheduler) to share data between blocks.




# The Problem

- Funding Source: Confidential high-throughput SDR communications system (multi-MIMO) that runs on a high-end, multi-core x86 processor (dozens) server with high-end memory configuration.
- They found that their throughput scales poorly as CPU count and memory are increased; they could not support additional real-time channels by introducing more hardware (within reason).
- After debugging and analysis, the problem was found to be memory bandwidth contention.
- Problem: Shared-memory system scales poorly with thread count.



# Potential Solutions

- Distribute block threads across disjoint memories and then use message passing between partitions.
- Use existing TCP/UDP blocks and break flowgraph into networked chain
  - No dynamic load balancing; bottleneck slows down the entire system.

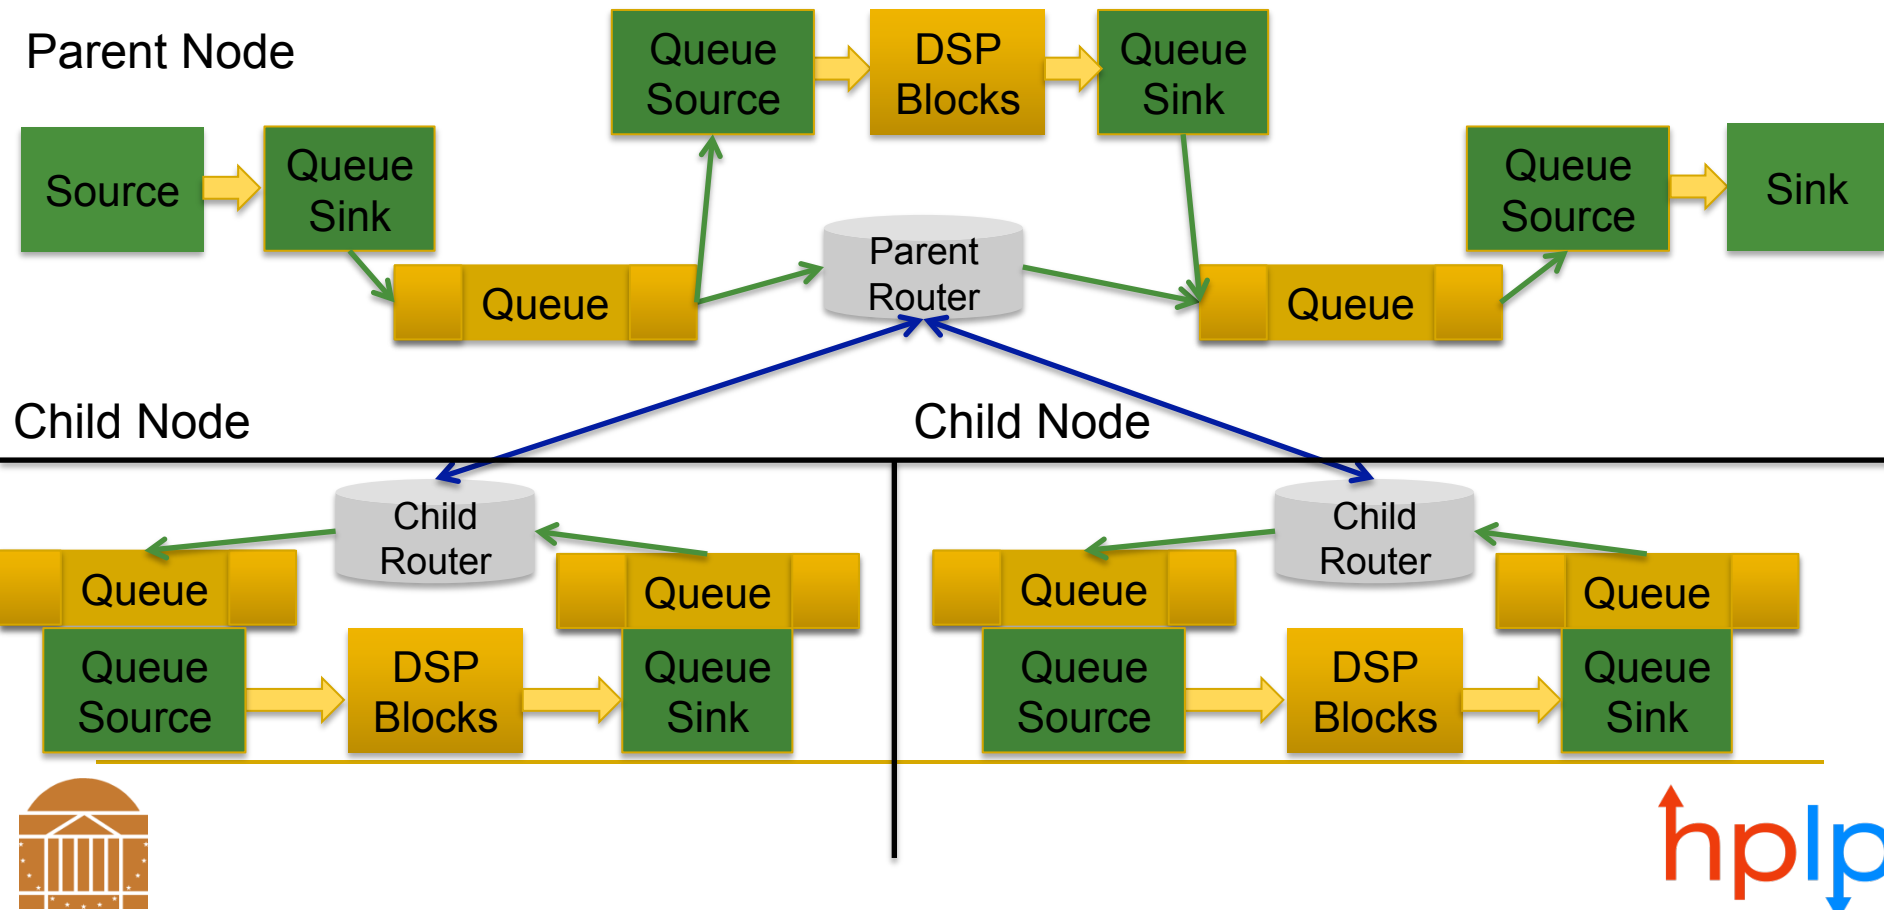


- Instead use dynamic load-balancing to distribute the workload across multiple nodes; with the nodes running redundant blocks.

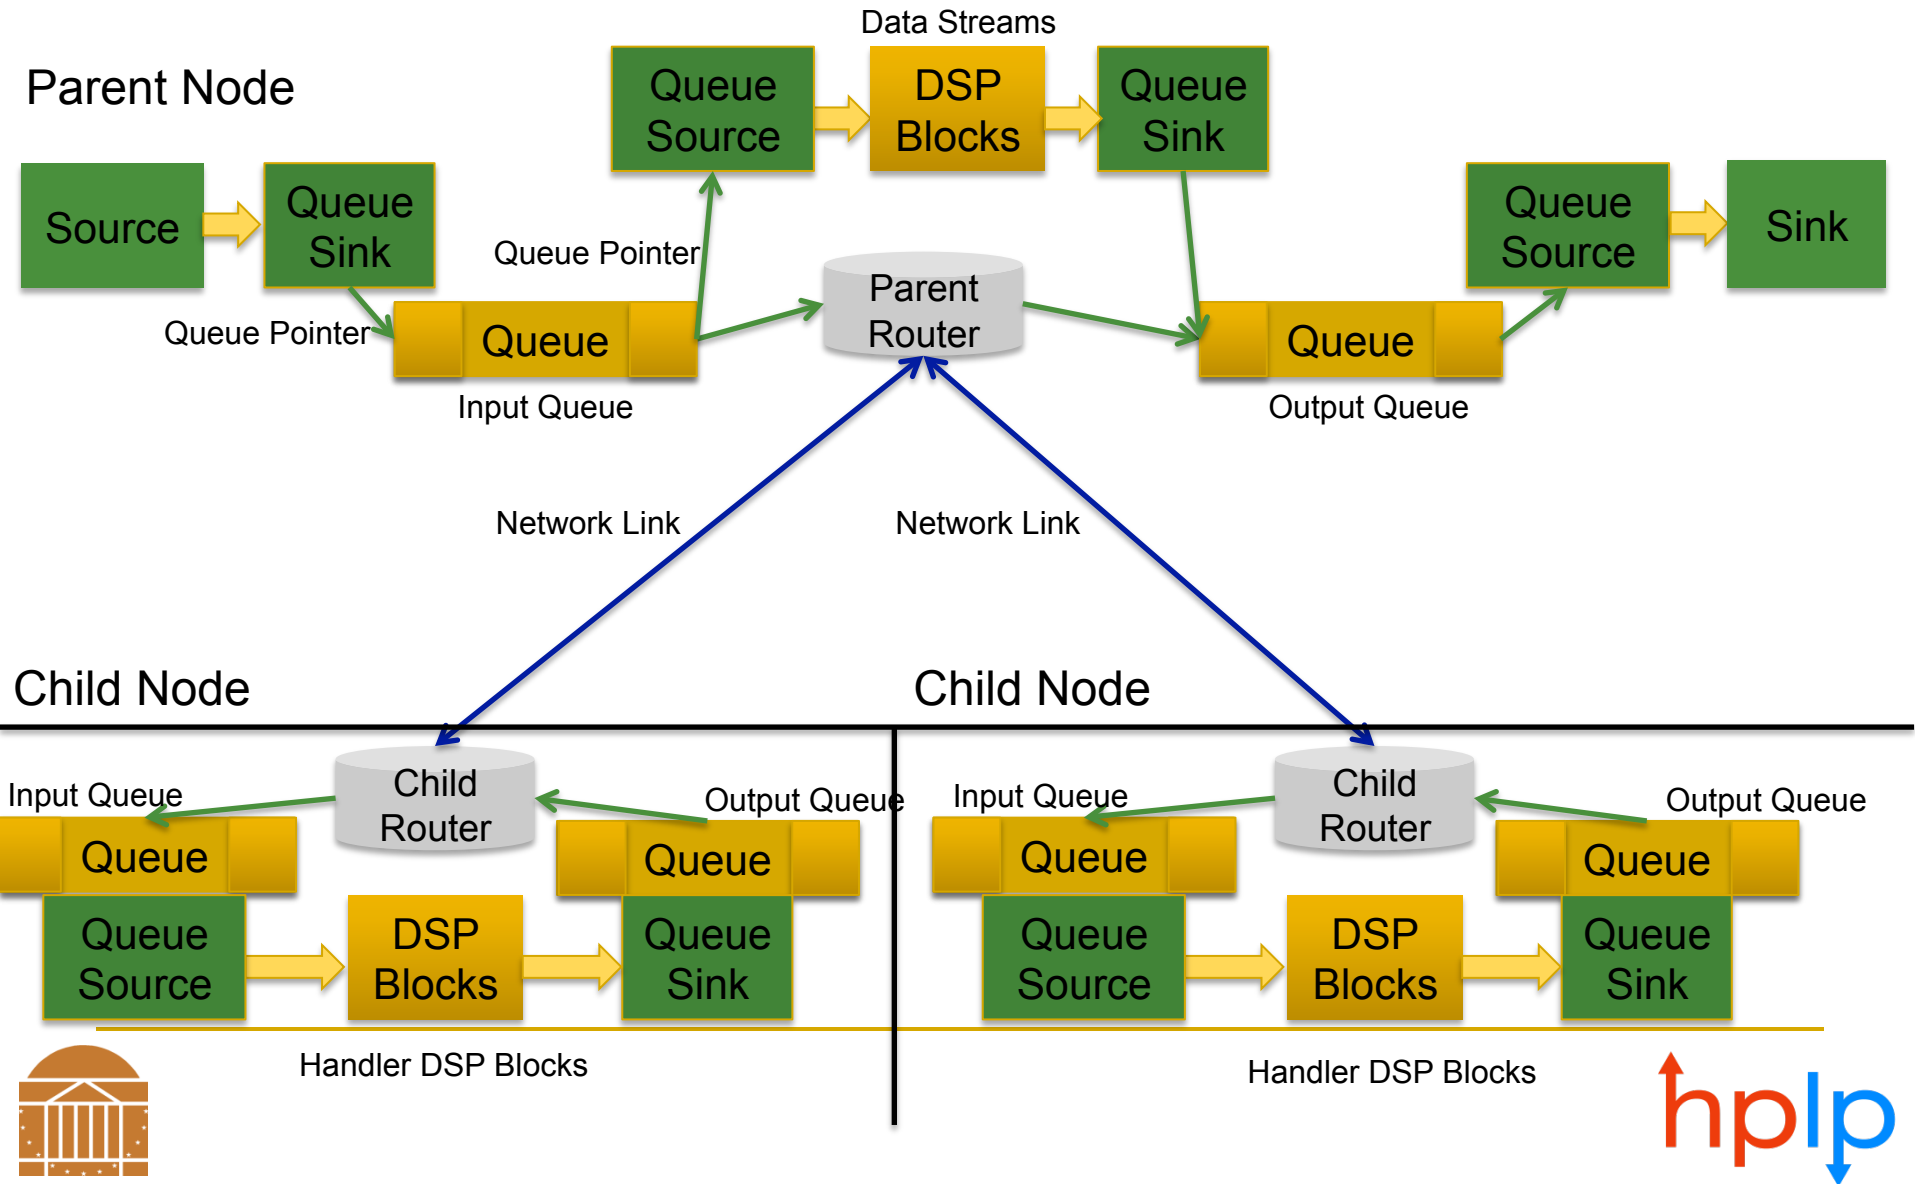


# Solution: GR-Router

- GR-Router: Distributed GNU Radio framework
- Distributes the DSP workload across multiple machines tied to a common communication network



# Terminology



# Parent Code

```
// Populate input queue
tb->connect(src, 0, throttle_0, 0);
tb->connect(throttle_0, 0, input_queue_sink, 0);

// Handler Code
tb->connect(input_queue_source, 0, decoder, 0);
tb->connect(decoder, 0, unpacked2packed, 0);
tb->connect(unpacked2packed, 0, output_queue_sink, 0);

//Order and sink results
tb->connect(output_queue_source, 0, throughput, 0);
tb->connect(throughput, 0, sink, 0);
tb->run();
```

# Child Code


```
// Handler code
tb->connect(input_queue_source, 0, decoder, 0);
tb->connect(decoder, 0, unpacked2packed, 0);
tb->connect(unpacked2packed, 0, output_queue_sink, 0);
tb->run();
```






# Blocks in GR-Router


- Queue Source and Sink Blocks



- Parent & Child Router Blocks




- Sync and Sink Throughput blocks



# Queue\_Sink Blocks

- **Queue Sink:** Converts stream to segments; pushes them onto queue
  - ❑ **Input:** Stream of data
  - ❑ **Output:** Pointer to queue
  - ❑ Supports reconstructing index
- **Segments:** Defined within the queue sink and source




Type 1 Msg.	type	index	size	Data[] : 768 floats	
Type 2 Msg.	type	index	size	Data[] : 50 bytes	
Type 3 Msg.	type	index	size	Data[] : 50 bytes	weight



# Queue\_Source Blocks

- **Queue Source:** Converts segments from a queue to streaming data
  - **Input:** Pointer to queue
  - **Output:** Stream of data
  - Supports ordering by index, and preserving index across stream with stream tags.

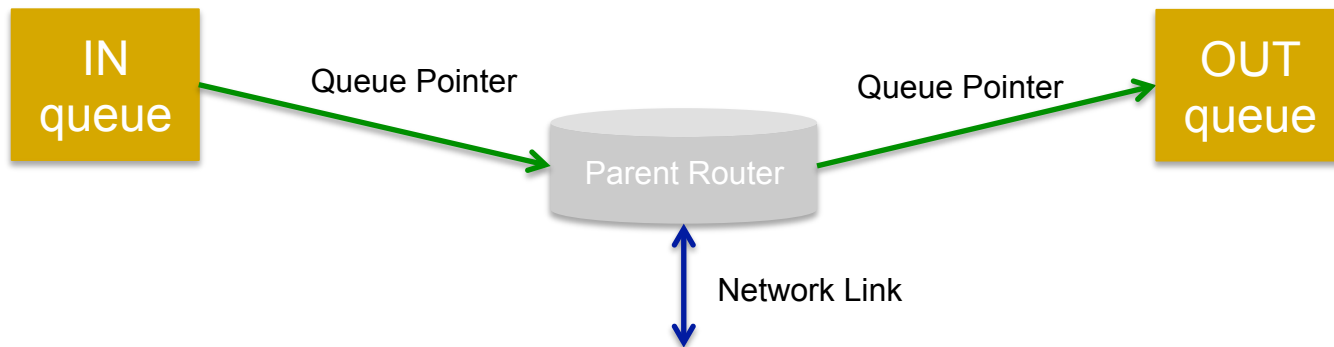


Type 1 Msg.	type	index	size	Data[] : 768 floats	
Type 2 Msg.	type	index	size	Data[] : 50 bytes	
Type 3 Msg.	type	index	size	Data[] : 50 bytes	weight




# Parent\_Router Block

- **Parent Router:** evenly distributes data segments among children
  - ❑ Keeps track of childrens' weights and distributes segments from the IN queue to balance the network.
  - ❑ Receives result segments from children (containing their 'weight')
  - ❑ **Input:** pointer from IN queue
  - ❑ **Output:** pointer to OUT queue
  - ❑ **Network Link:** Connection to Children.



# Child\_Router Block


- **Child Router:** computes segments and returns results to Parent with weight.
  - ❑ Keeps track of current 'busy-ness', and sends weight to Parent
  - ❑ **Input:** Pointer from OUT queue
  - ❑ **Output:** Pointer to IN queue
  - ❑ **Network Link:** Connection to parent.



# {Inline, Sink}\_Throughput Blocks


- **Inline Block:** In-line transparent block to calculate throughput

- ❑ **Input:** Previous block
- ❑ **Output:** Next block



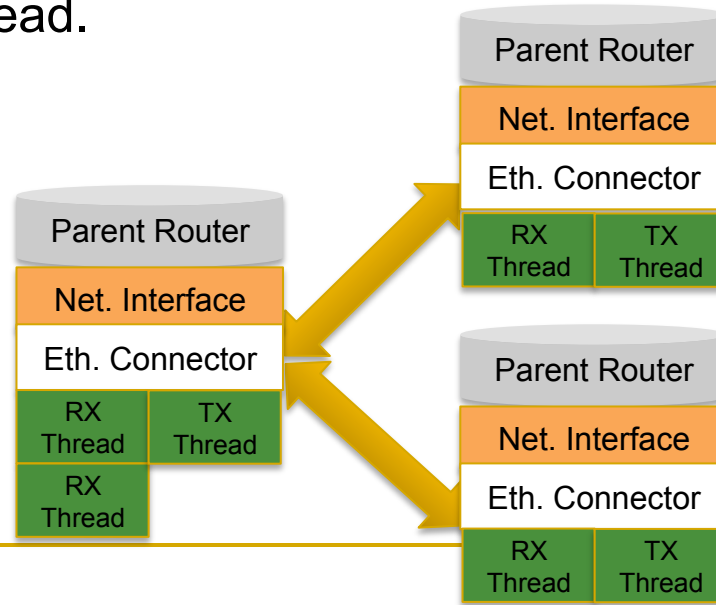
- **Sink Block:** Out-of-line sink block used to calculate throughput

- ❑ **Input:** Previous block
- ❑ **Output:** None



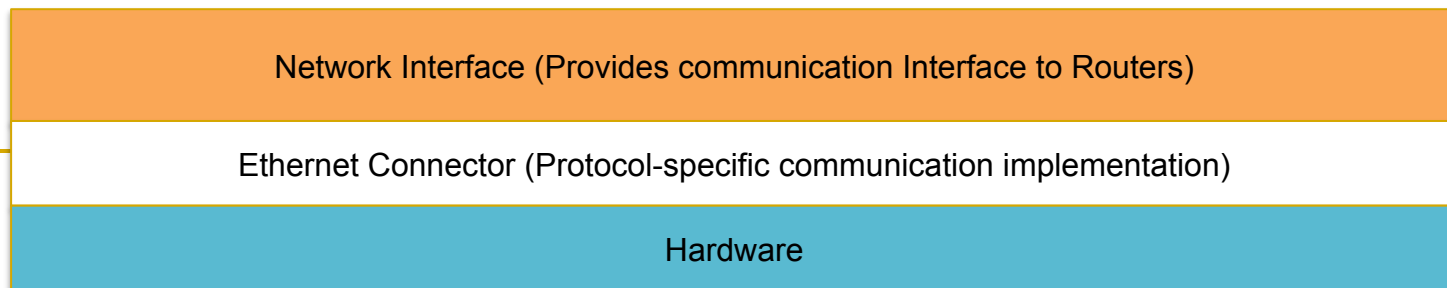

# Communications

- We used Ethernet for communication between machines.
  - GR-Router was designed to work with any available communication technologies.
  - For each node there is a receive(RX) thread per child router, and a single transmit(TX) thread.



# Network Interface

- Network Interface: Provides communication functions to Parent and Child Routers
  - **Connect**(char\* hostname): connects the graph of routers
    - Each child specifies it's parent, and each parent knows the number of children it connects to.
  - **Receive**(int index, char\* outbuf, int num\_items): receives segments from node at index
  - **Send**(int index, char\* msg, int num\_items): sends message to given node
  - Uses Connector for communications; is only meant as a uniform interface for the routers to abstract away the technology-specific implementation





# Ethernet Connector

- Ethernet Connector: Ethernet-specific networking class
- Provides the following functions:
  - ❑ `connect_to_parent(char* hostname, int port)`
  - ❑ `write_parent(char* msg, int size), read_parent(char* outbuf, int size)`
  - ❑ `connect_to_child(int index, int port)`
  - ❑ `write_child(int index, char* inbuf, int size), read_child(int index, char* outbuf, int size)`
  - ❑ `Stop()`
- In order to use a different technology (PCIe), this connector would need to be written for that technology

Ethernet Connector (Protocol-specific communication implementation)

Hardware (Ethernet, PCIe, etc)




# Example Implementation

- **GR-LDPC** : Manu T S wrote an implementation of a Low-Density Parity-Check Code encoder and decoder for GSoC 2013.
- Goal: Parallelize his decoder to increase net throughput.

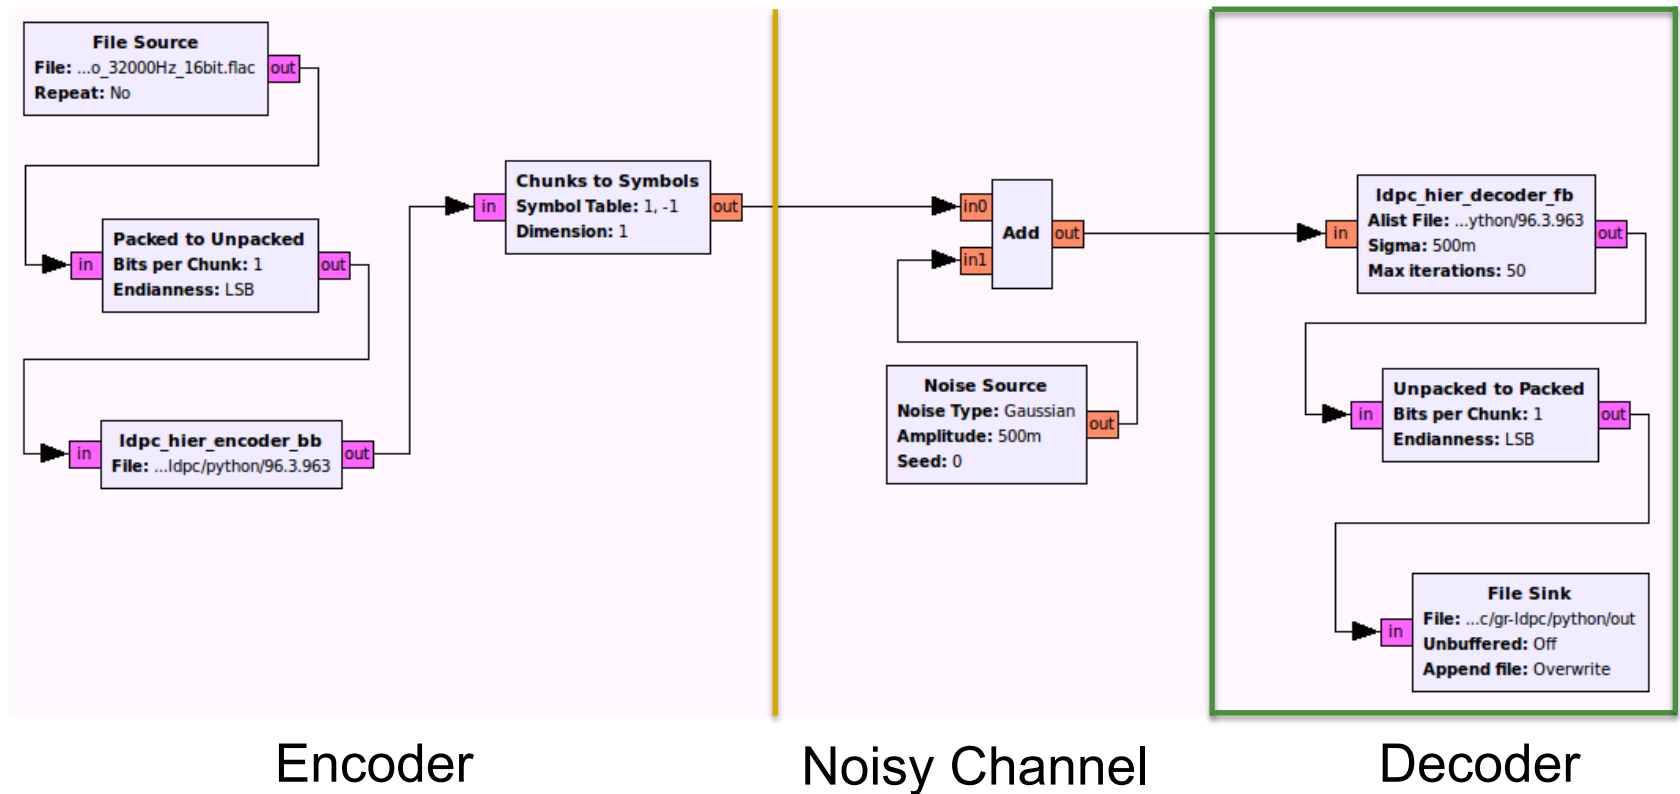


# LDPC Decoding

- **LDPC**: forward error correcting code using a soft decision decoder to 'guess and check' until a valid code word is found.
  - once found, the data can be calculated or a LUT can do the translation.
- It makes these decisions independently on the n-bit code word. Therefore LDPC is data-parallelizable. Multiple code-words can be decoded in parallel.



# Manu's LDPC Example

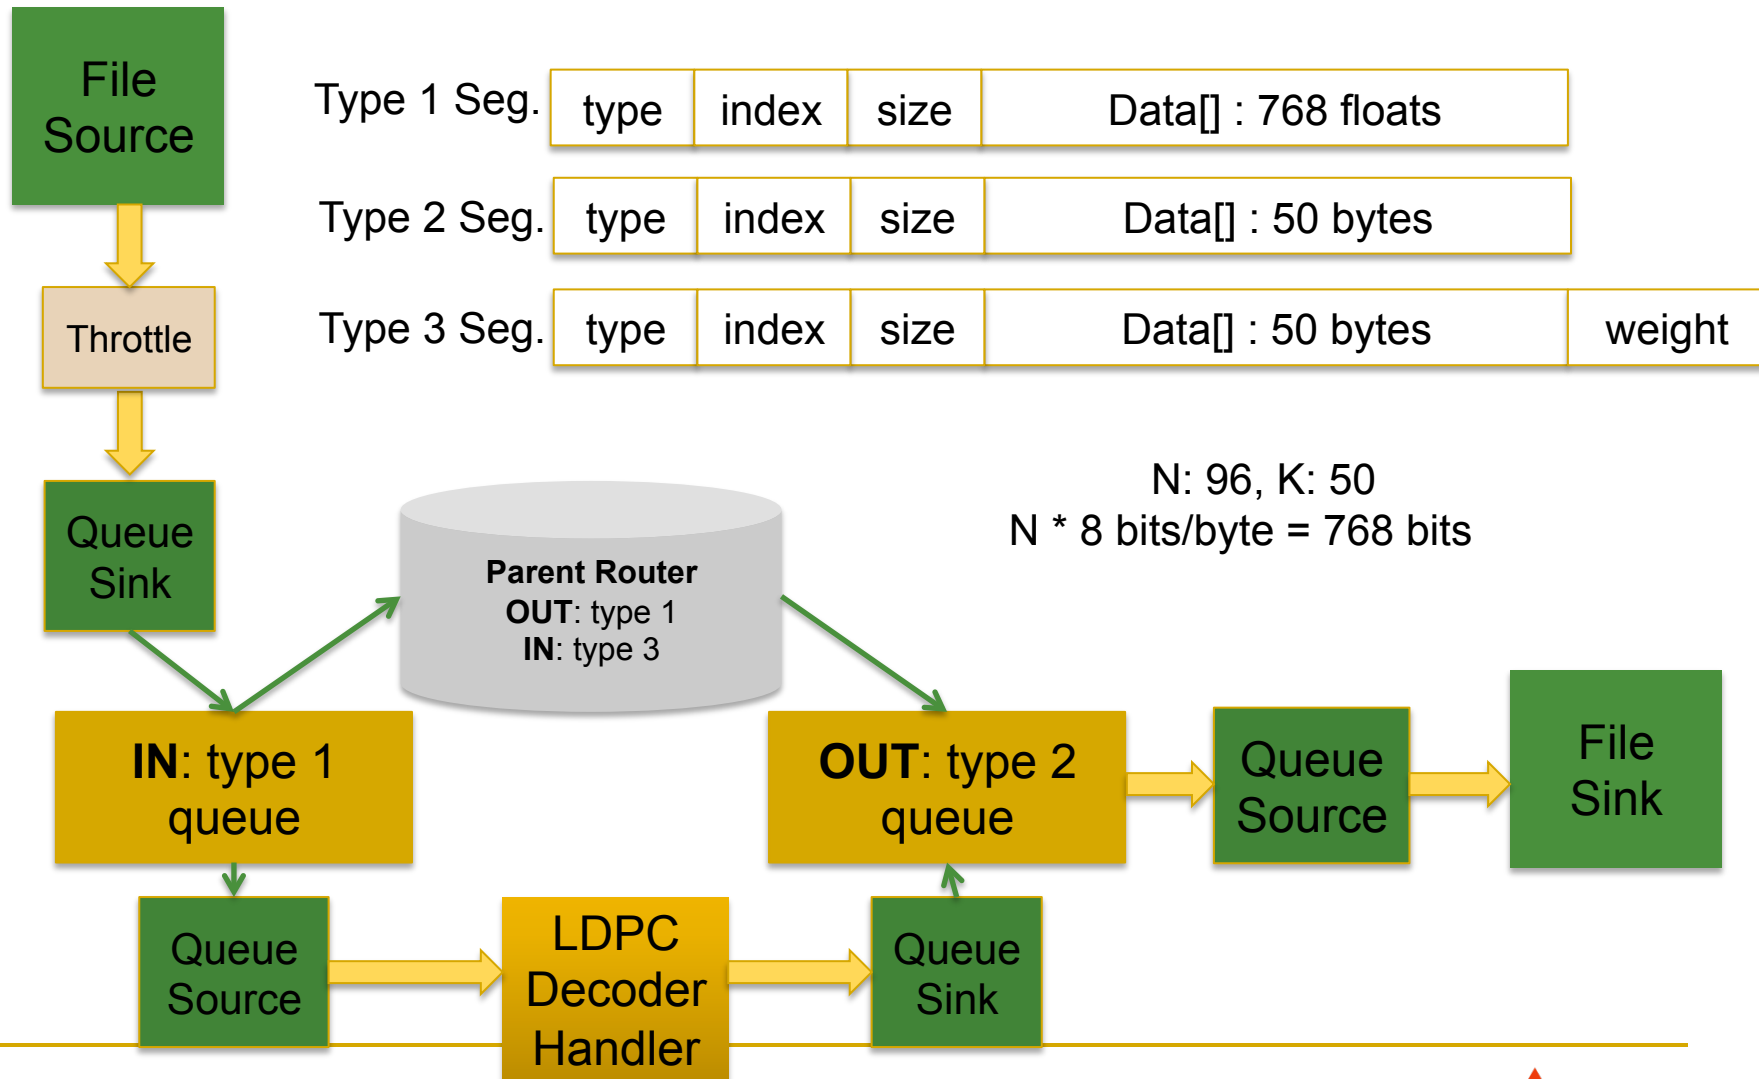


# Parallelization Strategy

- 1. Port Manu's code to C++
- 2. Create no-noise and 0.5 pk-pk LDPC-encoded files.
- 3. Decode each of the inputs on each node and determine maximum throughput that the node can maintain without GR-Router. (base case)
- 4. Create a parent application that serves as the root of the graph.
- 5. Create a child application for all nodes that connect to the root.
- 6. Add nodes and compare throughput results.



# LDPC Decoder Parent



# LDPC Decoder Child

Type 1 Msg. 

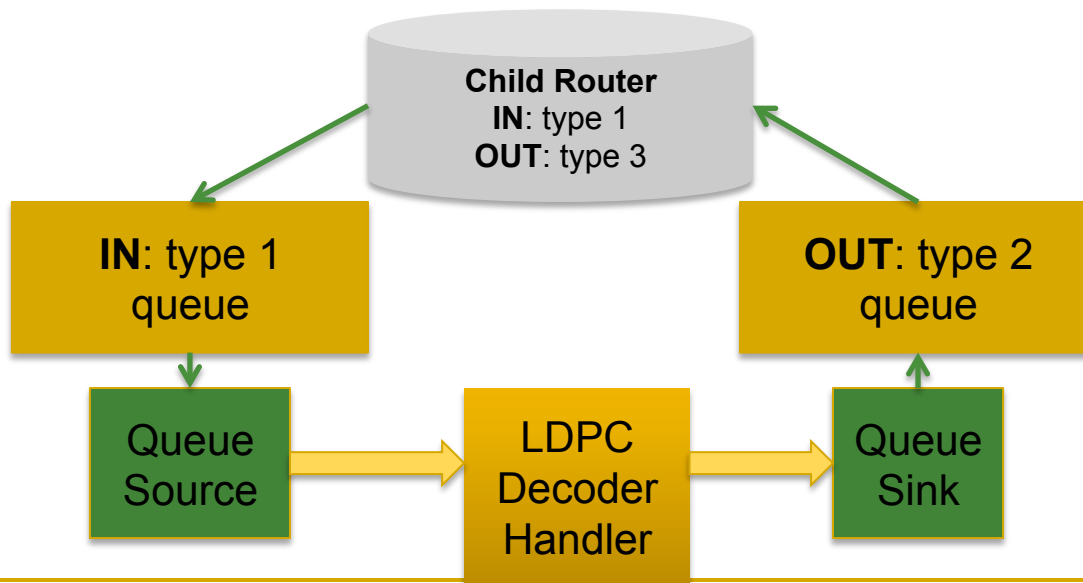
type	index	size	Data[] : 768 floats
------	-------	------	---------------------

Type 2 Msg. 

type	index	size	Data[] : 50 bytes
------	-------	------	-------------------

Type 3 Msg. 

type	index	size	Data[] : 50 bytes	weight
------	-------	------	-------------------	--------



N: 96, K: 50  
 $N * 8 \text{ bits/byte} = 768 \text{ bits}$



# Throughput Results

- Results:
  - Overhead: 20% reduced throughput on single node
  - Without noise: 20% reduced throughput when running on two nodes.
  - With 0.5 pk-pk noise: 1.8x throughput on two nodes!
- Currently running tests with 2-6 ethernet-connected machines for additional testing.





# Future Work

- Continue adding features to GR-Router
  - XML integration for easy reconfiguration
    - Define segments in an xml file, and have the `queue_{source, sink}` and routers read them.
- Implement redundancy and fault tolerance.
- Support additional network technologies.
- Support multi-layer GNU Radio systems.
- Support inter-child communication.
- Experiment with networked low-power platforms.



<http://www.iconarchive.com/show/adobe-cs4-icons-by-hopstarter/File-Adobe-Dreamweaver-XML-01-icon.html>



---

# Questions?

